

[30 points - 3 hours]

1) Linear Models, Linear Regression [5 points]

(a) Consider a classification problem.  $D = \{x_i, y_i\}_{i=1}^N$ . let  $x \in \mathbb{R}^d$   
 $y \in \{0, 1\}$ . Let there be  $n_1$  positive datapoints and  $n_0$  negative datapoints.

(i) Now  $\mu_1$  is the mean of the datapoints s.t  $y = 1$ . And  $\mu_0$  is the mean of features of negative datapoints.

Find  $\mu_0, \mu_1$ . (1)

(ii) Now you make a classifier  $h(x)$  that classifies the datapoints based on the proximity to the means  $\mu_0, \mu_1$ , i.e

$$\hat{y}_{\text{pred}} = \underset{i=0,1}{\operatorname{argmin}} \|x - \mu_i\|^2$$

Find the decision boundary of this model. (2)

(b) Say, you want to predict the number of runs a team will score using 100 features. But you only have 80 training examples. Argue why a linear regression model will not have a well-defined solution. (2)

[Hint:  $y = Xw$  - Think in terms of the solution of  $w$  and the rank of  $X$ ]

2) Discriminant Functions [5 points]

Consider the discriminant functions for discriminating between 2 classes.

$$g_1(x) = P(C_1 | x), g_2(x) = P(C_2 | x); \text{s.t.}$$

$$\text{Given } x | C_1 \sim N \left( \begin{bmatrix} 1 \\ 2 \end{bmatrix}, \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix} \right)$$

$$\text{Given } x | c_1 \sim N \left( \begin{bmatrix} 1 \\ 2 \end{bmatrix}, \begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix} \right) \quad \det P(c_1) = 0.6$$

$$x | c_2 \sim N \left( \begin{bmatrix} 0 \\ -1 \end{bmatrix}, \begin{bmatrix} 2 & 0 \\ 0 & 2 \end{bmatrix} \right) \quad p(c_2) = 0.4$$

(i) Find the decision boundary :  $g_1(x) = g_2(x)$ . (3)

(ii) Hence plot the distribution of data (using contours) and the decision boundary using python. (2)

[Use the same jupyter notebook for all the subsequent questions].

3) Support Vector Machines ; Sequential Minimal Optimization [7 points]

SVM objective (Recall.)

$$\underset{\alpha}{\text{maximize.}} \quad \ell(\alpha) = \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha_i \alpha_j y_i y_j K(x_i, x_j)$$

$$\text{s.t. } \sum_{i=1}^n \alpha_i y_i = 0 \quad y = \{-1, 1\}$$

$$0 \leq \alpha_i \leq C$$

Now, we know to optimize this objective we use coordinate ascent with SMO algorithm; where only 2 variables are optimized in 1 step. Here let us only consider 1 step of the optimization as : ( $\alpha_1, \alpha_2$  is being optimized)

Since we need to optimize only over  $\alpha_1, \alpha_2$  (others: constant)

$$\alpha_1^*, \alpha_2^* = \underset{\alpha_1, \alpha_2}{\operatorname{argmax}} \quad w_0 + w_1 \alpha_1 + w_2 \alpha_2 - \left\{ w_3 \alpha_1 \alpha_2 + w_4 \alpha_1^2 + w_5 \alpha_2^2 \right\}$$

$$\text{where } \alpha_1 y_1 + \alpha_2 y_2 = t$$

$$0 \leq \alpha_1, \alpha_2 \leq C$$

Note: you will need this detail for coding - (only for giving a valid input.)  
in matrix form

$$\underset{\alpha_1, \alpha_2}{\text{eqn max}} \quad w_0 + [w_1, w_2] \begin{bmatrix} \alpha_1 \\ \alpha_2 \end{bmatrix} = \begin{bmatrix} w_1 & w_2 \\ w_3/2 & w_4 \end{bmatrix} \begin{bmatrix} \alpha_1 \\ \alpha_2 \end{bmatrix}$$

$$\begin{bmatrix} w_1 & w_3/2 \\ w_3/2 & w_4 \end{bmatrix} \stackrel{(\text{psd})}{\geq 0} \iff (\text{According to SVM this should be psd})$$

- (i) find  $\alpha_1^*, \alpha_2^*$ . [ Hint : Session 19 - full SMO was taught ]  
[ 5 ] - [ In terms of the constants  $w_0, w_1, w_2, w_3, y_1, y_2, t, C$  ]
- (ii) Hence implement it using Python. [ 2 ]

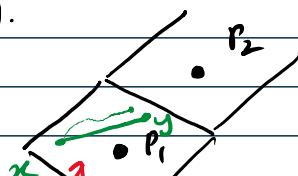
Inputs:  $(w_0 \text{ to } w_5)$ ;  $y_1, y_2 \in \{-1, 1\}$ ;  $t, C$

Outputs:  $\alpha_1^*, \alpha_2^*$

#### 4) Nearest Neighbour Algorithm [ 3 points ]

Consider a 1-nearest neighbour algorithm. Show that it divides the space into Dirichlet/Voronoi tessellation. The partitioning of the space is into 'convex polygons'; s.t. each polygon has exactly 1 generating point :

i.e. There is only one point that points inside the polygon are close to. (Ignore points at the boundary of the polygons).





all pts here belong to  $P_i$ .

In other words:  $\theta x + (1-\theta)y \in \text{Polygmn}$  if  $x, y \in \text{Polygmn}$

[Hint:  $\theta x + (1-\theta)y \in \text{Polygmn}$  if  $x, y \in \text{Polygmn}$   
 $\theta \in [0, 1]$ ] - all pts connecting  $x$  and  $y$  have to  
 be contained in the polygon.

Note: you can assume L2-distance.

[Hint: The cosine rule may be handy!]

## 5) Softmax Regression and MLP [10 points]

### Dataset

Use the MNIST dataset [handwritten digits  $\rightarrow (0-9)$ ];  
 Use sklearn to import it as:

```
import numpy as np
from sklearn.datasets import fetch_openml
```

X,y=fetch\_openml("mnist\_784",version=1,return\_X\_y=True,as\_frame=False)

### MLP

(i) Divide the dataset into 40k (train), 10k (CV set), 20k (test)  
 [8.7 they have equal number of 0's, 1's, 2's ... 9's]

- This is called a stratified split. [1]

(ii) Now you can use the backprop code I shared. Modify it  
 to make it predict a multiclass classification [0-9].  
 (10 way)

The last layer has to have 10 nodes.

Test it with 1, 2, and 3 hidden layers. Select the  
 best model using the CV accuracy. Report the Accuracy,  
 and contingency table on test set. (10x10 table)

Contingency table

Contingency table

	Predicted				
Actual		$c_1$	$c_2$	$c_3$	$c_4$
$c_1$	/ / /				
$c_2$		/ / /			
$c_3$			/ / /		
$c_4$				/ / /	

$c_{ij}$ : The number of points in class  $i$  predicted as  $c_j$

(iii) Logistic Regression [6 marks]

Hence; divide the dataset so you have to perform.

binary classification - between odd and even numbers.

Train set size : 60K

Test set size : 10K

$0/p \rightarrow 0$  (even number)  
 $1$  (odd number) [Science implement a linear regression model]

Report: (i) Accuracy

(ii) Contingency table ( $2 \times 2$ )

(iii) Precision (P) :=  $\frac{(\# \text{True } +ve)}{(\# \text{True } +ve) + (\# \text{False } +ve)}$

[% - of your +ve predictions that were correct]  $\xrightarrow{\text{you said you, actually}}$

(iv) Recall (R) :=  $\frac{(\# \text{True } +ve)}{(\# \text{True } +ve) + (\# \text{False } -ve)}$  no

[% - the class data pts you missed]

(v) F1 Score :=  $\frac{2PR}{P+R}$  [3 points]

Predicted

		True (1)	False (0)
		True +ve	False -ve
True (1)	/ / /		
False		False -ve	True
False			...

: (Traditional Contingency table)

meaning:

(True / False) ( +ve / -ve)

	False	True	(True / False)	(True / False)
False (0)	+ve	-ve	↓	↓
			correct prediction	(predicted class 1)
			(wrong prediction)	(predicted class 0)

Note : Packages allowed (Python)

- (i) numpy
- (ii) pandas
- (iii) matplotlib
- (iv) sklearn (only to import data - no predefined functions must be used)

### Implementation Tips

- 1) Images ( $n \times n$ ) → flatten it to make  $n^2 \times 1$  before using a MLP. { i.e. convert the matrix to a vector }
- 2) Don't forget to divide the image matrix by 255 - so that the image is scaled between 0, 1 . It helps in stable training.